



IIT School of Applied Technology

ILLINOIS INSTITUTE OF TECHNOLOGY

information technology & management

526 Data Warehousing

March 30, 2016

Week 11 Presentation

Week 11 Topic: **SQL Optimization** for DW


Partial vs. Full Range Scan

- We will cover
 - Partial/Full Range Scan Concept
 - Partial vs. Full Range Scan & Examples
 - BI Applications Demo
 - Assignment 02 Specifications


SQL Optimization for DW

Partial vs. Full Range Scan

- Access subset of the entire result before pushing out the first delivery buffer
- Improve response time



Full Range Scan



Partial Range Scan

SQL Optimization for DW


Partial Range Scan

- Using Index to replace SORT operation
- Using Index to replace MAX operation
- Access index file only without accessing the data file
- Using EXISTS
- Using ROWNUM
- Using Stored Function


SQL Optimization for DW

Using Index to Replace SORT

```
SELECT *  
  FROM PRODUCT  
 WHERE YMD = '20151023'  
    AND ITEM LIKE 'AB%'  
 ORDER BY YMD, ITEM;
```



```
SELECT *  
  FROM PRODUCT  
 WHERE YMD = '20151023'  
    AND ITEM LIKE 'AB%';
```



SQL Optimization for DW

Partial Range Scan: Replacing SORT

```
SELECT ORDDATE, CUSTNO
  FROM ORDER1T
 WHERE ORDDATE between
        '940101' and '941130'
 ORDER BY ORDDATE DESC
```

5.2 sec

```
21200  SORT ORDER BY
21200  TABLE ACCESS BY ROWID ORDER1T
21201  INDEX RANGE SCAN ORD_ORDDATE
```

```
SELECT /*+ INDEX_DESC(A orddate) */
      ORDDATE, CUSTNO
  FROM ORDER1T A
 WHERE ORDDATE between
        '940101' and '941130'
```

0.01 sec

```
20 INDEX RANGE SCAN DESCENDING ORDDATE
```

```
SELECT ORDDATE, CUSTNO
  FROM ORDER1T
 WHERE ORDDATE LIKE '7%'
 ORDER BY ORDDATE DESC
```

12.5 sec

```
42000  SORT ORDER BY
42000  TABLE ACCESS BY ROWID ORDER1T
42001  INDEX RANGE SCAN ORD_DEPT
```

```
SELECT /*+ INDEX_DESC(A orddate) */
      ORDDATE, CUSTNO
  FROM ORDER1T A
 WHERE ORDDATE LIKE '7%'
        AND ORDDATE <= '991231'
```


0.02 sec

```
20 INDEX RANGE SCAN DESCENDING ORDDATE
```

SQL Optimization for DW

Partial Range Scan: Access Index Only


```
SELECT DEPT, SUM(QTY)
FROM PRODUCT
WHERE DEPT LIKE '12%'
GROUP BY DEPT;
```



Index:
DEPT

Table

```
SELECT DEPT, SUM(QTY)
FROM PRODUCT
WHERE DEPT LIKE '12%'
GROUP BY DEPT;
```




Index:
DEPT + QTY


SQL Optimization for DW

Partial Range Scan: Handling MAX

```
SELECT MAX (SEQ) + 1  
FROM PRODUCT  
WHERE DEPT = '12300';
```



```
SELECT /*+ INDEX_DESC (A IX_DEPT_SEQ) */  
      A.SEQ+1  
FROM PRODUCT A  
WHERE A.DEPT LIKE '12300'  
      AND ROWNUM = 1;
```



SQL Optimization for DW

Partial Range Scan: SQL Examples

```
SELECT MAX(ORDDATE)
FROM ORDER1T
WHERE ORDDEPT = '430'
AND STATUS = '30'
```

2.53 sec

```
1      SORT AGGREGATE
2892      TABLE ACCESS BY ROWID ORDER1T
15230      INDEX RANGE SCAN DEPT_DATE
```

```
SELECT /*+ INDEX_DESC(A dept_date) */
      ORDDATE
FROM ORDER1T A
WHERE ORDDEPT = '430' AND STATUS='30'
AND ROWNUM = 1
```

0.01 sec

```
1 COUNT STOPKEY
1      TABLE ACCESS BY ROWID ORDER1T
2      INDEX RANGE SCAN DESCENDING DEPT_DATE
```

```
SELECT TYPE, COUNT(*)
FROM ORDER2T
WHERE ITEM LIKE 'HJ%'
GROUP BY TYPE
```

10.3 sec

```
20      SORT GROUP BY
36631      TABLE ACCESS BY ROWID ORDER2T
36631      INDEX RANGE SCAN ITEM_STATUS
```

```
SELECT STATUS, COUNT(*)
FROM ORDER2T
WHERE ITEM LIKE 'HJ%'
GROUP BY STATUS
```


2.5 sec

```
20      SORT GROUP BY
36631      INDEX RANGE SCAN ITEM_STATUS
```


SQL Optimization for DW

Partial Range Scan: EXISTS

```
SELECT COUNT(*) INTO :CNT
  FROM ITEM_TAB
 WHERE DEPT = '101'
    AND SEQ > 100
....
IF CNT > 0 ...
....
```



```
SELECT 1 INTO :CNT FROM DUAL
 WHERE EXISTS
   (SELECT 'X'
    FROM ITEM_TAB
   WHERE DEPT = '101'
      AND SEQ > 100)
....
IF CNT > 0
....
```



SQL Optimization for DW



Partial Range Scan: ROWNUM

```
SELECT COUNT(*) INTO :CNT
FROM ITEM_TAB
WHERE DEPT = '101'
AND SEQ > 100
```

```
....
IF CNT > 0
....
```

```
SELECT 1 INTO :CNT
FROM ITEM_TAB
WHERE DEPT = '101'
AND SEQ > 100
AND ROWNUM = 1
```

```
....
IF CNT > 0
....
```



SQL Optimization for DW

Partial Range Scan: 1:M JOIN

```
SELECT x.CUST_NO, x.ADDR, x.NAME, ...
  FROM CUST x, REQT y
 WHERE x.CUST_NO = y.CUST_NO
    AND x.CUST_STAT IN ('A', 'C', 'F')
    AND y.UN_PAY > 0
 GROUP BY x.CUST_NO
HAVING SUM(y.UN_PAY) BETWEEN :VAL1 and :VAL2
```

Throughput

Full Range Scan

Response Time

Partial Range Scan

```
SELECT x.CUST_NO, x.ADDR, x.NAME, ....
  FROM CUST x
 WHERE CUST_STAT in ('A', 'C', 'F')
    AND EXISTS
      (SELECT 'X'
        FROM REQT y
       WHERE y.CUST_NO = x.CUST_NO
          AND UN_PAY > 0
      GROUP BY x.CUST_NO
      HAVING SUM(y.UN_PAY)
        BETWEEN :VAL1 and :VAL2
      )
```

SQL Optimization for DW

Partial Range Scan: 1:M JOIN (cont'd)

```

SELECT x.CUST_NO, SUM(y.UN_PAY) AS UN_PAY, .
  FROM CUST x, REQT y
 WHERE x.CUST_NO = y.CUST_NO
    AND x.CUST_STAT IN ('A', 'C', 'F')
    AND y.UN_PAY > 0
  GROUP BY x.CUST_NO
HAVING SUM(y.UN_PAY) BETWEEN :VAL1 and :VAL2

```

Throughput

Full Range Scan

```

CREATE or REPLACE FUNCTION unpay_sum
  (v_custno in varchar2)
  return number is
    sum_unpay number;
begin
  ....
  select sum(un_pay)
    into sum_unpay
  from reqt
 where cust_no = v_custno
    and un_pay > 0;
  ....
  return sum_unpay;
end unpay_sum ;

```

Response Time

Partial Range Scan

```

SELECT cust_no, addr, un_pay, ....
  FROM (

```

```

  SELECT cust_no, addr,
    unpay_sum(cust_no)
    AS un_pay,
    ....

```

```

  FROM cust
  WHERE cust_stat IN
    ('A', 'C', 'F')

```

```

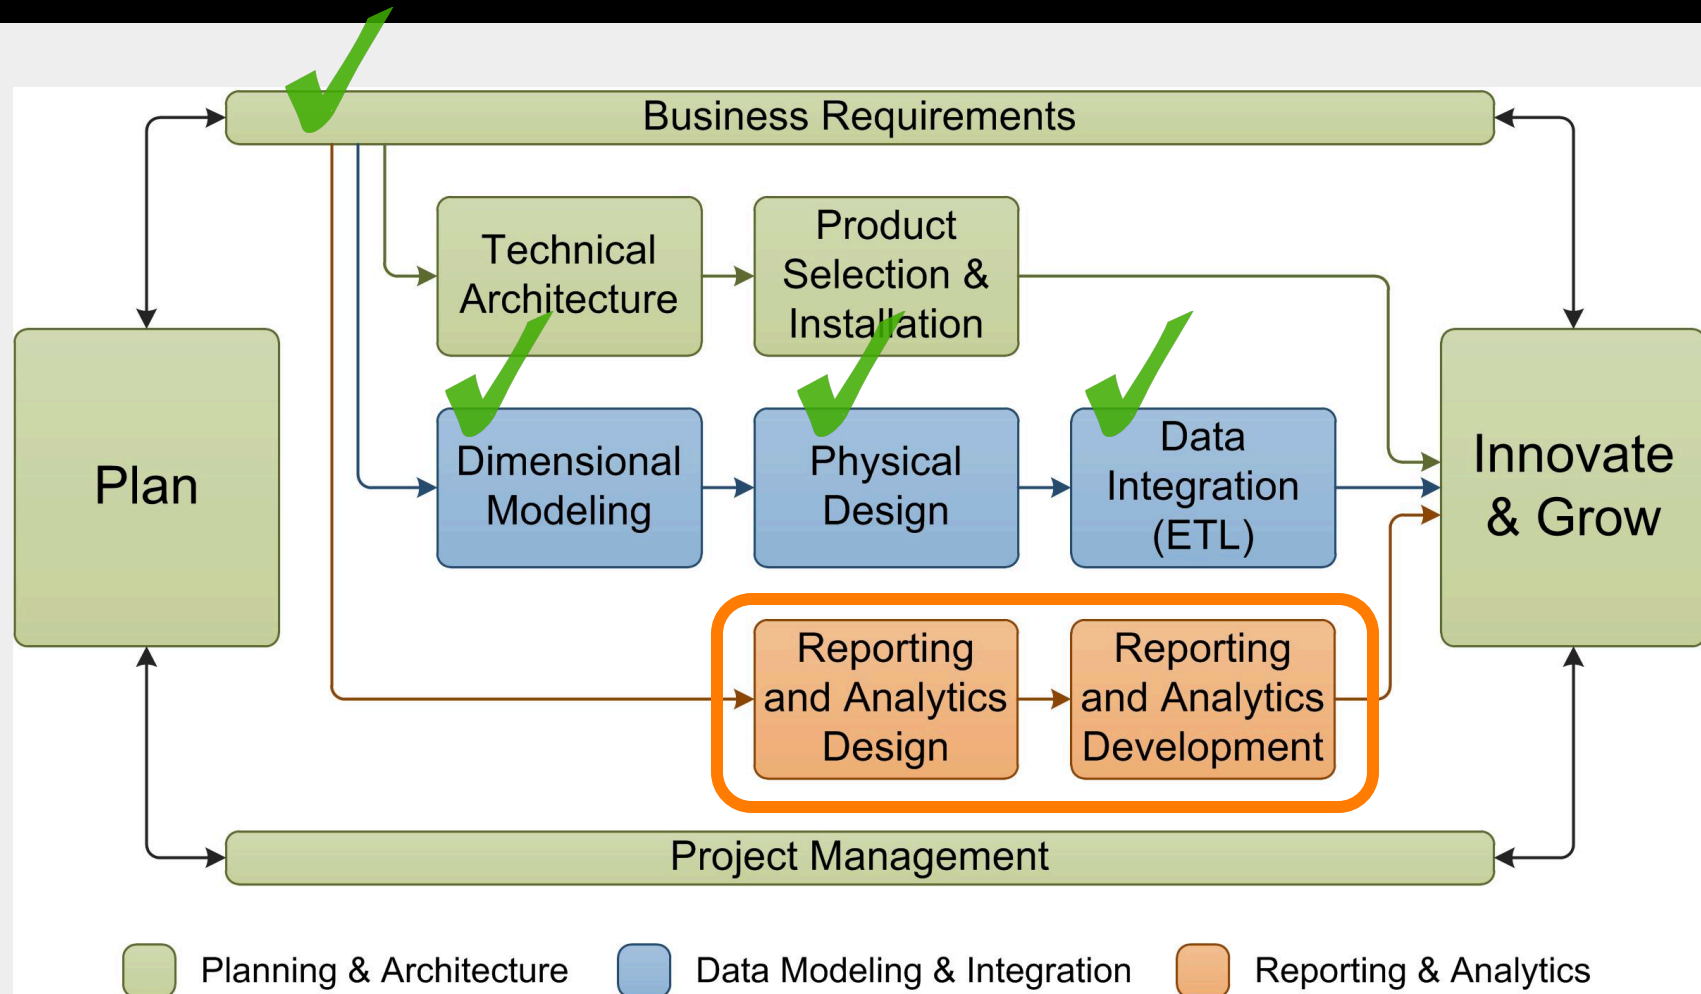
)
WHERE un_pay BETWEEN :VAL1 AND :VAL2

```

Week 11 Class Exercises

- BI Applications Demo
 - Metadata Editor
 - Metadata Workbench
 - Saiku Analytics (Mondrian OLAP)
 - Pentaho BI Suite
 - Tableau (<http://www.tableau.com/academic/students>)
- The Last Home Work (Group Home work)
 - Saiku Analytics Installation
- Assignment 02 Preview
 - Proof Of Concept Demo
 - Other Details

Kimball Lifecycle Approach Revisited





DW/BI Deliverables Revisited

- Dimensional models
 - Relational **star schemas**
 - Multidimensional online analytical processing (**OLAP**) cubes
- Business intelligence applications

Star Schema Versus OLAP Cubes

- At a logical level, there is no difference
- It is a matter of physical database implementation.
- Star schema is implemented in a relational database and is queried through SQL
- OLAP Cubes (multidimensional databases) are implemented for extreme performance and are queried through MDX.



Star Schema vs. OLAP Cubes Revisited

- The star schema can store large amounts of detailed data.
- OLAP Cubes provide higher performance with pre-calculated summary data.
- In general, OLAP cubes are populated from the star schema

- Kimball focuses on the star schema rather than the OLAP cubes
- Star schema usually has 15 dimensions
- OLAP usually has 8-10 dimensions

Assignment 02 Specification (1/3)

1. Dimensional Model

■ Dimensions

- Minimum 1 Role Playing Dimension
- 1 Junk Dimension (DIM_ETHNICITY)
- Minimum 1 Bridge Table
- Suggested Total Number of Dimensions: 5 ~ 8

■ Measures

- GRE percentiles (Quant., Analytics, and Verbal)
- Numeric flags on admission status

PROG_ACTN ▼	IS_APPLIED ▼	IS_ADMITTED ▼	IS_ACCEPTED ▼
APPL	1	0	0
DENY	1	0	0
WAPP	1	1	0
MATR	1	1	1

Assignment 02 Specification (2/3)

2. ETL Implementation

▪ Dimensions

- Implement minimum 5 dimensions
- No bridge table implementation necessary

▪ Measures

- GRE percentiles (Quant., Analytics, and Verbal)
- Numeric flags on admission status

3. Reporting and Analytics

- Minimum 5 findings on the following measures using BI application(s)
 - Selectivity and Yield Ratio
 - GRE percentiles
- BI application options
 - Saiku Analytics, Tableau, SQL Drilling down, etc.

Assignment 02 Specification (3/3)

4. Project Deliverables

- Dimensional Model (in any format)
- Fully functional ETL codes
- Documentations
 - Setup instructions & DDL Scripts
 - ETL execution instructions
 - Minimum 5 findings from the Reporting & Analytics

5. Extra Credit

- Group presentation on April 27
- Must be notified by April 20th (FCFS)
- The credit can be applied to all but the Final