



IIT School of Applied Technology
ILLINOIS INSTITUTE OF TECHNOLOGY
information technology & management

526 Data Warehousing

April 6 & 13, 2016

Week 12 & 13 Presentation

Week 12 Topic: SQL Optimization for DW

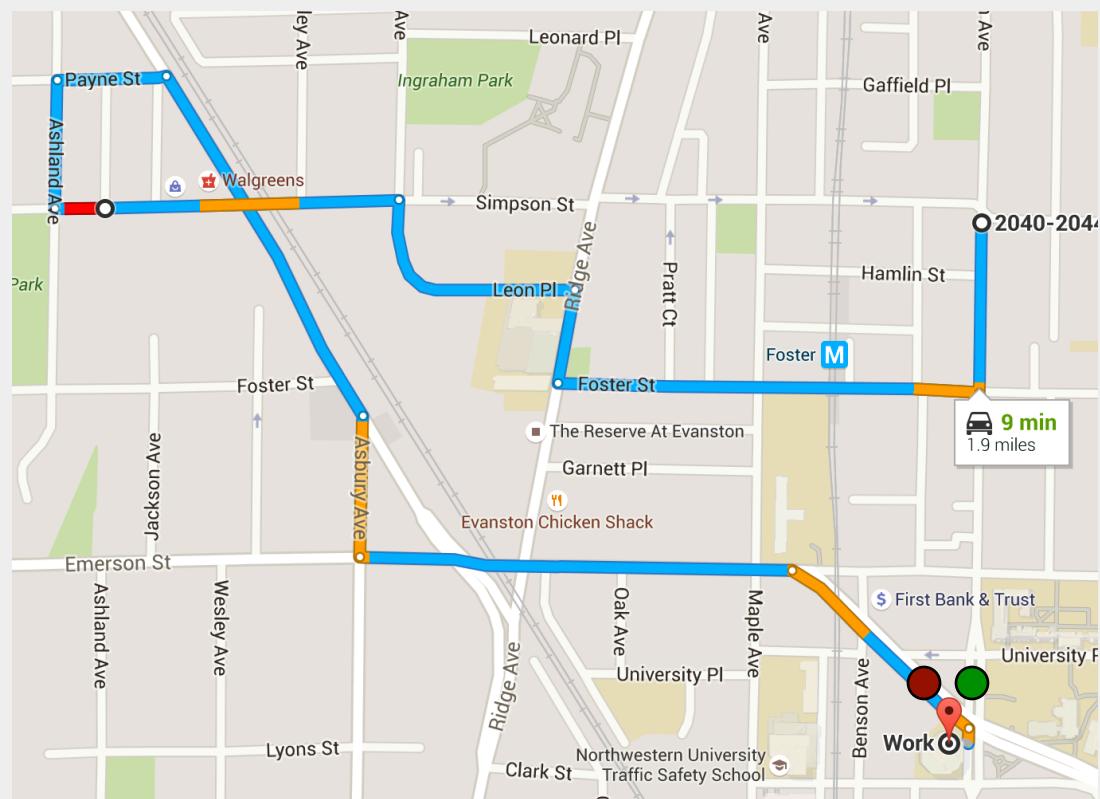
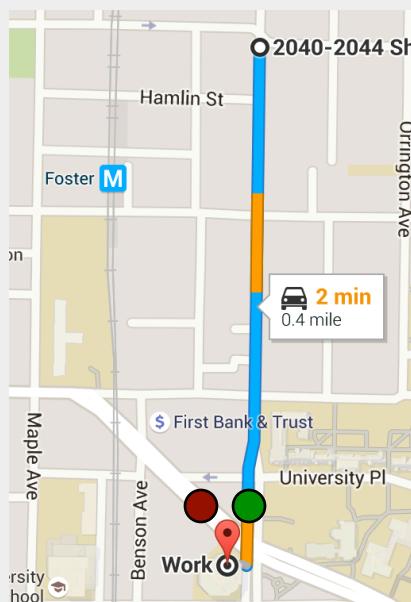
Introduction of Index

- We will cover
 - Unusable Index
 - External Index Suppression
 - Internal Index Suppression
 - Index with NOT Operator
 - Index with NULL & NOT NULL

Week 12 Topic: SQL Optimization for DW

Optimizer is Not Perfect

- Bad Index
- Good Index



SQL Optimization for DW Unusable Index

- Index can not be leveraged when...

Index column suppressed

```
SELECT *
  FROM DEPT
 WHERE SUBSTR(DNAME,1,3) = 'ABC'
```

Using NOT operator

```
SELECT *
  FROM EMP
 WHERE JOB <> 'SALES'
```

Using NULL or NOT NULL

```
SELECT *
  FROM EMP
 WHERE ENAME IS NOT NULL
```

Another index is preferred

```
SELECT *
  FROM EMP
 WHERE JOB LIKE 'AB%'
   AND EMPNO = '7890'
```

SQL Optimization for DW

External Index Suppression

```
SELECT *  
  FROM EMP  
 WHERE SUBSTR(DNAME,1,3) = 'ABC'
```

```
SELECT *  
  FROM EMP  
 WHERE DNAME LIKE 'ABC%'
```

```
SELECT *  
  FROM EMP  
 WHERE SAL * 12 = 12000000
```

```
SELECT *  
  FROM EMP  
 WHERE SAL = 12000000 / 12
```

```
SELECT *  
  FROM EMP  
 WHERE TO_CHAR(HIREDATE, 'YYMMDD')  
       = '20140101'
```

```
SELECT *  
  FROM EMP  
 WHERE HIREDATE =  
       TO_DATE('20140101','YYYYMMDD')
```

```
SELECT *  
  FROM EMP  
 WHERE NVL(COMM,0) < 100
```

?

SQL Optimization for DW (cont'd)

External Index Suppression

```
SELECT *  
  FROM EMP  
 WHERE EMPNO BETWEEN 100 AND 200  
   AND NVL(JOB,'X') = 'CLERK'
```

```
SELECT *  
  FROM EMP  
 WHERE EMPNO BETWEEN 100 AND 200  
   AND JOB = 'CLERK'
```

```
SELECT *  
  FROM EMP  
 WHERE DEPTNO||JOB = '10SALESMAN'
```

```
SELECT *  
  FROM EMP  
 WHERE DEPTNO = '10'  
   AND JOB = 'SALSMAN'
```

Intentional Suppression

```
SELECT *  
  FROM EMP  
 WHERE JOB = 'MANAGER'
```

```
SELECT *  
  FROM EMP  
 WHERE RTRIM(JOB) = 'MANAGER'
```

```
SELECT *  
  FROM EMP  
 WHERE EMPNO = 8978
```

```
SELECT *  
  FROM EMP  
 WHERE RTRIM(EMPNO) = 8978
```

SQL Optimization for DW (cont'd)

External Index Suppression

```
SELECT CUSTNO, INVDATE  
  FROM INVOICE  
 WHERE CUSTNO = 'DN02'  
   AND STATUS = '90'
```

TABLE ACCESS BY ROWID INVOICE
AND-EQUAL
INDEX RANGE SCAN INV_STATUS
INDEX RANGE SCAN INV_CUSTNO

4 rows,
0.51 sec

```
SELECT CUSTNO, INVDATE  
  FROM INVOICE  
 WHERE CUSTNO = 'DN02'  
   AND RTRIM(STATUS) = '90'
```

TABLE ACCESS BY ROWID INVOICE
INDEX RANGE SCAN INV_CUSTNO

4 rows,
0.03 sec

```
SELECT CUSTNO, INVDATE  
  FROM INVOICE  
 WHERE CUSTNO LIKE 'DN%'  
   AND STATUS LIKE '9%'
```

rows,
sec

```
SELECT CUSTNO, INVDATE  
  FROM INVOICE  
 WHERE CUSTNO LIKE 'DN%'  
   AND RTRIM(STATUS) LIKE '9%'
```

rows,
sec

SQL Optimization for DW (cont'd)

External Index Suppression

```
SELECT X.CUSTNO, INVDATE, CUSTNAME  
      FROM SALES01 X, SALES02 Y  
     WHERE X.SALENO = Y.SALENO  
       AND X.SALEDEPT = '710'  
       AND Y.SALEDAT E LIKE '201411%'
```

```
SELECT X.CUSTNO, INVDATE, CUSTNAME  
      FROM SALES01 X, SALES02 Y  
     WHERE X.SALENO = Y.SALENO  
       AND RTRIM(X.SALEDEPT) = '710'  
       AND Y.SALEDAT E LIKE '201411%'
```

```
SELECT X.ORDNO, ORDDATE, ITEM  
      FROM ORDER01 X, ORDER02 Y  
     WHERE X.ORDNO = Y.ORDNO  
       AND X.ORDDATE LIKE '201411%'  
       AND Y.ORDDEPT = '710'  
    ORDER BY ORDDATE
```

```
SELECT X.ORDNO, ORDDATE, ITEM  
      FROM ORDER01 X, ORDER02 Y  
     WHERE RTRIM(X.ORDNO) = Y.ORDNO  
       AND X.ORDDATE LIKE '201411%'  
       AND Y.ORDDEPT = '710'
```

SQL Optimization for DW Internal Index Suppression

```
SELECT *  
  FROM ATABLE  
 WHERE CHA = 10
```

```
SELECT *  
  FROM ATABLE  
 WHERE TO_NUMBER(CHA) = 10
```

```
SELECT *  
  FROM ATABLE  
 WHERE NUM LIKE '201410%'
```

```
SELECT *  
  FROM ATABLE  
 WHERE TO_CHAR(NUM) LIKE '201410%'
```

```
SELECT *  
  FROM ATABLE  
 WHERE DAT = '01-JAN-2014'
```

```
SELECT *  
  FROM ATABLE  
 WHERE DAT = TO_DATE('01-JAN-2014')
```

SQL Optimization for DW (cont'd)

Internal Index Suppression

```
SELECT *  
  FROM ATABLE  
 WHERE VAR = 10
```

```
SELECT *  
  FROM ATABLE  
 WHERE TO_NUMBER(VAR) = 10
```

```
SELECT *  
  FROM ATABLE  
 WHERE NUM = CHA
```

```
SELECT *  
  FROM ATABLE  
 WHERE NUM = TO_CHAR(NUM)
```

```
SELECT *  
  FROM ATABLE  
 WHERE DAT = '01-JAN-2014'
```

```
SELECT *  
  FROM ATABLE  
 WHERE DAT = TO_DATE('01-JAN-2014')
```

```
SELECT *  
  FROM ATABLE  
 WHERE DAT = CHA
```

```
SELECT *  
  FROM ATABLE  
 WHERE DAT = TO_DATE(CHA)
```

SQL Optimization for DW (cont'd)

Internal Index Suppression

```
SELECT SUM(UNCOST)  
  FROM INVOICE  
 WHERE STATUS = 90
```

1 row,
28.5 sec

TABLE ACCESS FULL INVOICE

```
SELECT SUM(UNCOST)  
  FROM INVOICE  
 WHERE STATUS = '90'
```

1 row,
0.15 sec

SORT AGGREGATE
TABLE ACCESS BY ROWID INVOICE
INDEX RANGE SCAN INV_STATUS

```
SELECT INVNO, CUSTNO, UNCOST  
  FROM INVOICE  
 WHERE CFMDEPT LIKE '71%'
```

rows,
sec

Number

TABLE ACCESS FULL INVOICE

```
SELECT ORDNO, INVNO, STATUS  
  FROM ORDER1T X, INVOICE Y  
 WHERE X.CUSTNO = Y.CUSTNO  
   AND X.ORDDEPT = Y.CFMDEPT  
   AND Y.INVDATE LIKE '201411%'
```

rows,
71 sec

NESTED LOOPS
TABLE ACCESS FULL ORDER1T
TABLE ACCESS BY ROWID INVOICE
INDEX RANGE SCAN INV_CFMDEPT

SQL Optimization for DW Index with NOT Operator

```
SELECT 'Not found !' INTO :COL1
  FROM EMP
 WHERE EMPNO <> '1234'
```

```
SELECT 'OK' INTO :COL1
  FROM DUAL
 WHERE NOT EXISTS (SELECT '' FROM EMP
                      WHERE EMPNO = '1234')
```

```
SELECT *
  FROM EMP
 WHERE ENAME LIKE 'JA%'
   AND JOB <> 'SALES'
```

```
SELECT *
  FROM EMP a
 WHERE a.ENAME LIKE 'JA%'
   AND NOT EXISTS (SELECT '' FROM EMP b
                      WHERE a.ENAME = b.ENAME
                        AND b.JOB = 'SALES')
```

```
SELECT *
  FROM EMP
 WHERE ENAME LIKE 'JA%'
MINUS
SELECT *
  FROM EMP b
 WHERE b.JOB = 'SALES'
```

SQL Optimization for DW Index with NULL & NOT NULL

```
SELECT *  
FROM EMP  
WHERE ENAME IS NOT NULL
```

```
SELECT *  
FROM EMP  
WHERE ENAME > ''
```

```
SELECT *  
FROM EMP  
WHERE COMM IS NOT NULL
```

```
SELECT *  
FROM EMP  
WHERE COMM > 0
```

```
SELECT *  
FROM EMP  
WHERE COMM IS NULL
```

?

SQL Optimization for DW

Optimizer Decision Factors

RANKING

```
SELECT *  
  FROM EMP  
 WHERE ENAME LIKE 'AB%'  
   AND EMPNO = '7890'
```

Use EMPNO Index Only

Avoid INDEX Merge

```
SELECT *  
  FROM EMP  
 WHERE ENAME LIKE 'AB%'  
   AND JOB LIKE 'SA%'
```

User either ENAME or
JOB index,
Or Full scan

Low COST

```
SELECT *  
  FROM EMP  
 WHERE EMPNO > '10'
```

FULL Table Scan

HINT

```
SELECT /*+ INDEX(EMP JOB_IDX) */ *  
  FROM EMP  
 WHERE ENAME LIKE 'AB%'  
   AND JOB LIKE 'SA%'
```

Use JOB Index Only

SQL Optimization for DW Optimizer: Cost vs. Rule Based

Rule Based

INDEX merge
(and_equal)

EMP Index

Index created recently

Cost Based

INDEX merge
(and_equal) or one
Index

EMP or ENAME Index
(depends on stats)

EMP or ENAME Index
or full scan (depends on
stats)

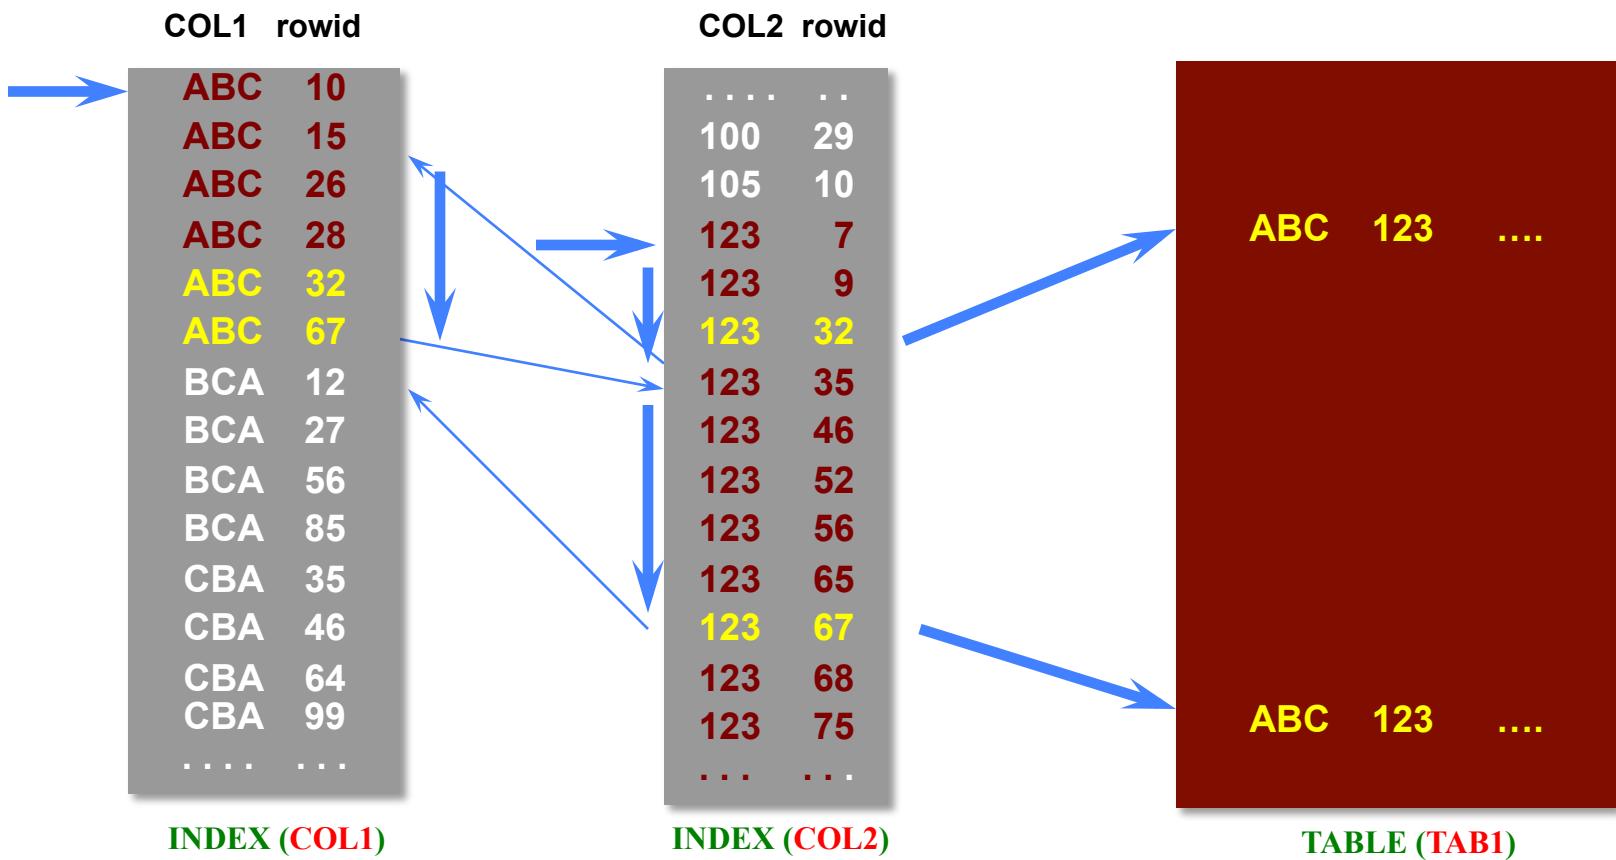
```
SELECT *  
FROM EMP  
WHERE JOB = 'SALESMAN'  
AND EMPNO = '7890'
```

```
SELECT *  
FROM EMP  
WHERE ENAME LIKE 'AB%'  
AND EMPNO = '7890'
```

```
SELECT *  
FROM EMP  
WHERE ENAME LIKE 'AB%'  
AND JOB LIKE 'SA%'
```

SQL Optimization for DW Index Merge

```
SELECT COL1, COL2, .....  
      FROM TAB1  
 WHERE COL1 = 'ABC' AND COL2 = '123' ;
```



SQL Optimization for DW

Index Merge vs. Composite Index

Index
Merge

```
SELECT COL1, COL2, .....  
      FROM TAB1  
     WHERE COL1 = 'ABC' AND COL2 = '123' ;
```

Composite
Index

COL1 rowid

ABC	10
ABC	15
ABC	26
ABC	28
ABC	32
ABC	67
BCA	12
BCA	27
BCA	56
BCA	85
CBA	35
CBA	46
CBA	64
CBA	99
....	...

COL2 rowid

....	...
100	29
105	10
123	7
123	9
123	32
123	35
123	46
123	52
123	56
123	65
123	67
123	68
123	75
....	...

INDEX (COL1)

INDEX (COL2)

COL1 COL2 rowid

ABC	55	10
ABC	67	15
ABC	89	26
ABC	100	29
ABC	123	32
ABC	123	67
ABC	180	76
BCA	100	12
BCA	100	27
BCA	123	56
BCA	123	85
CBA	85	35
CBA	123	46
CBA	214	64
....

INDEX (COL1 + COL2)

ITMD 526

SQL Optimization for DW

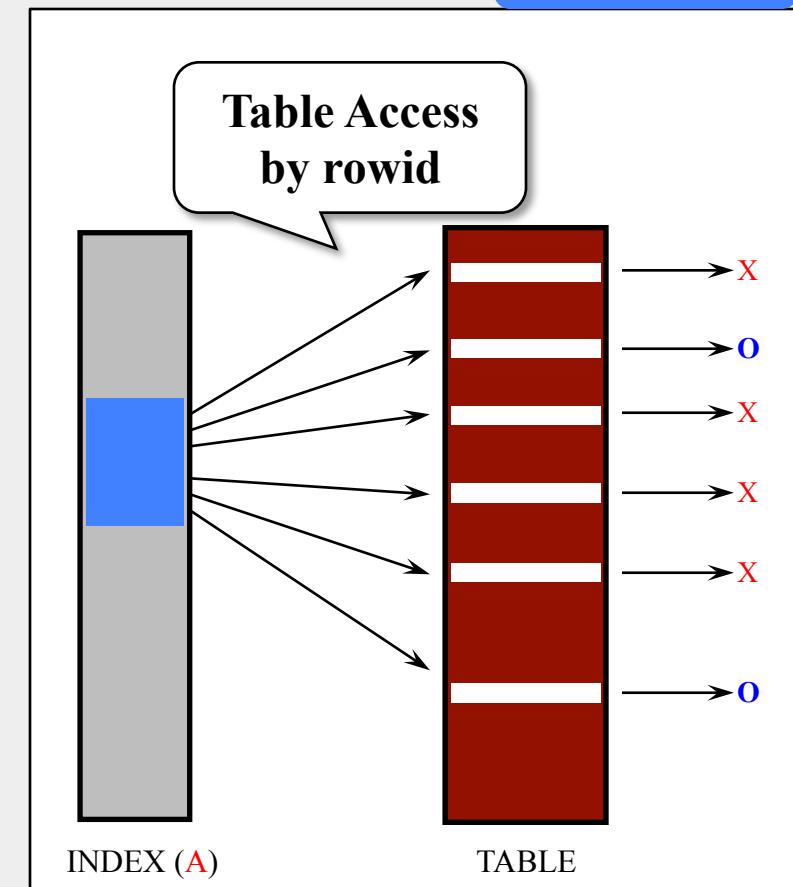
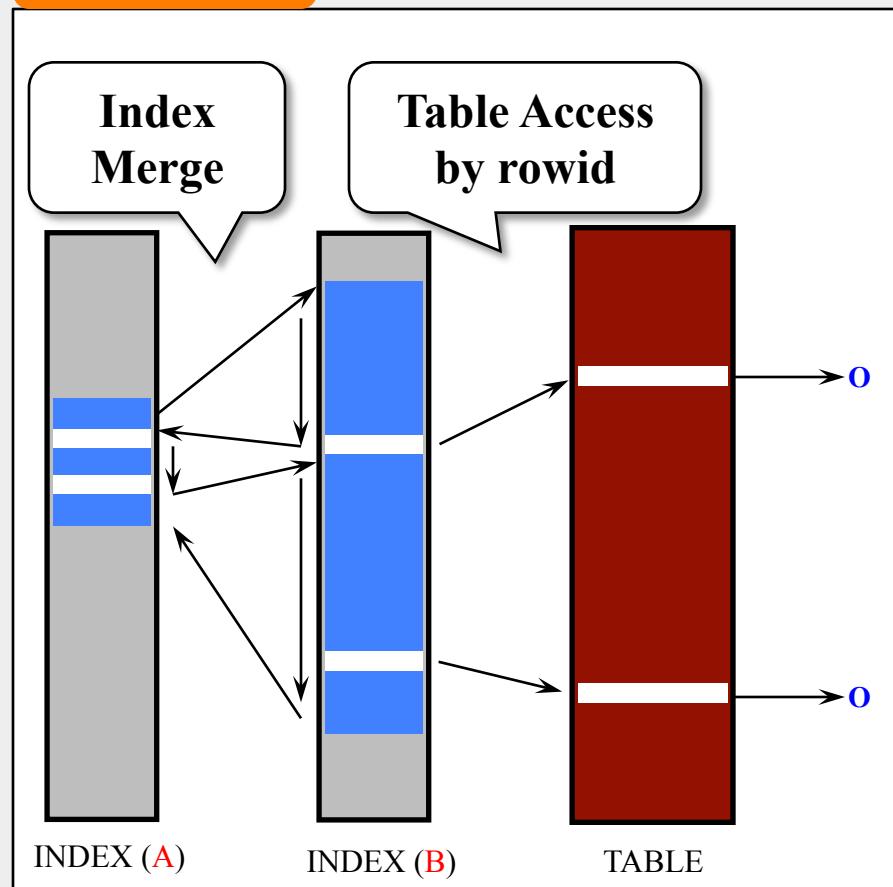
Index Merge - Suboptimal

ITMD - 526

```
SELECT COL1, COL2, ....  
  FROM TAB1  
 WHERE A = 'ABC' AND B LIKE '12%';
```

Suboptimal

Optimal



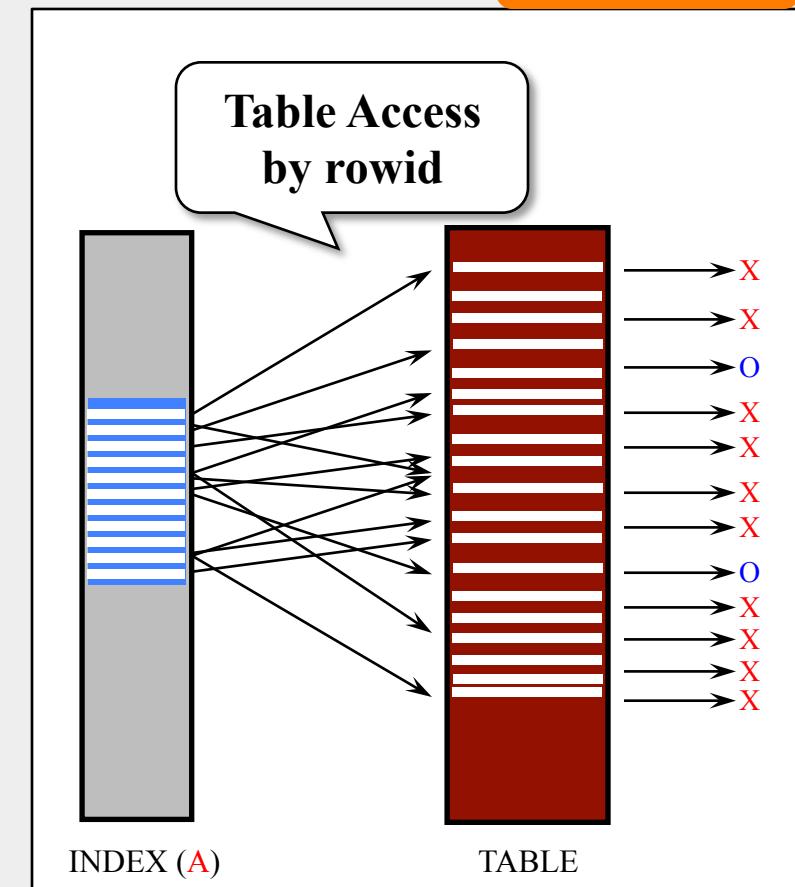
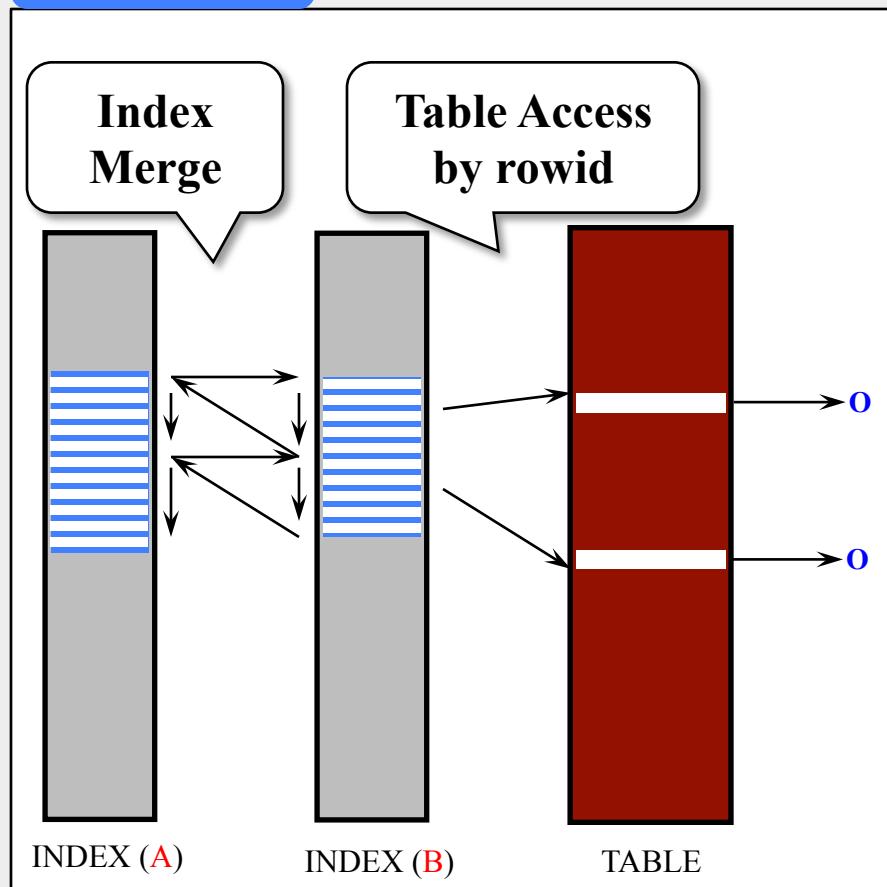
SQL Optimization for DW

Index Merge – Optimal

Optimal

```
SELECT COL1, COL2  
  FROM TAB1  
 WHERE A = 'ABC' AND B = '123';
```

Suboptimal



ITMD - 526

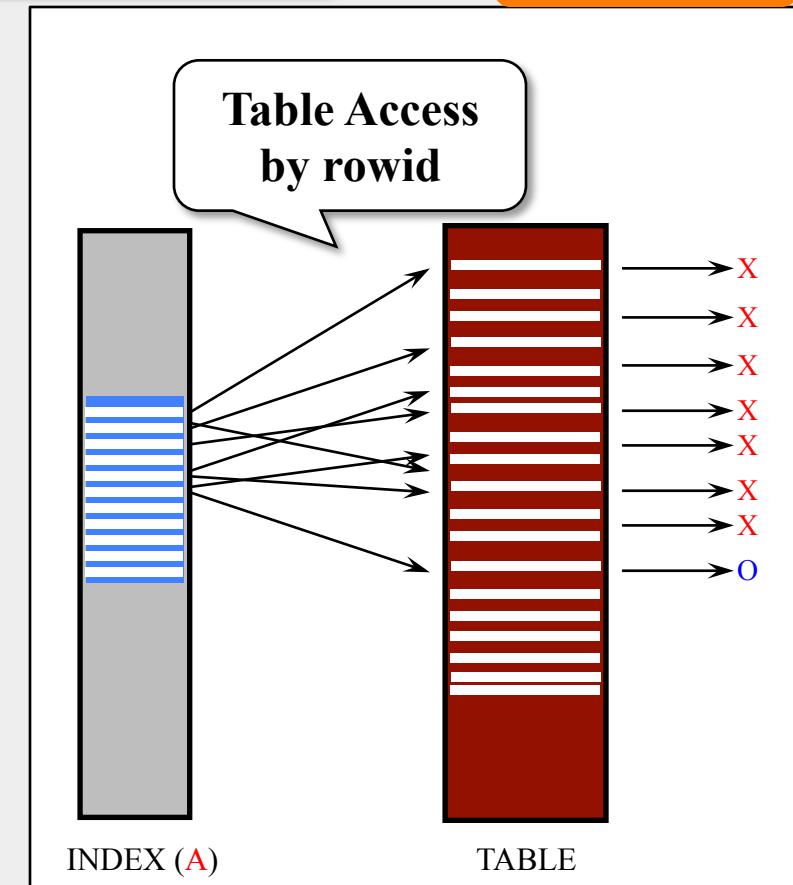
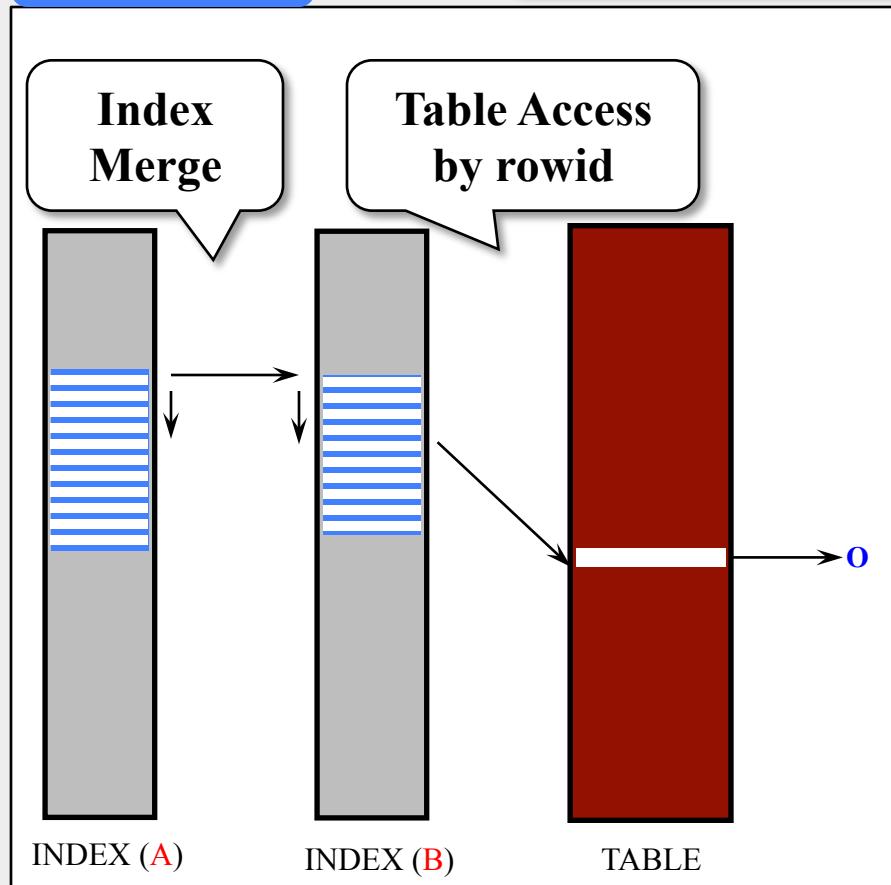
SQL Optimization for DW

Index Merge – Optimal (cont'd)

Optimal

```
SELECT 'X' FROM DUAL  
WHERE EXISTS (  
    SELECT 'X' FROM TAB1  
    WHERE A = 'ABC' AND B = '123');
```

Suboptimal



ITMD - 526

SQL Optimization for DW Index Merge Examples

```
SELECT CUSTNO, INVDATE  
      FROM INVOICE  
     WHERE CUSTNO = 'DN02'  
           AND INVDATE = '941003'
```

1 row,
0.02 sec

```
TABLE ACCESS BY ROWID INVOICE  
AND-EQUAL  
INDEX RANGE SCAN CH_CUSTNO  
INDEX RANGE SCAN CH_INVDATE
```

```
SELECT /*+ AND_EQUAL(A ch_custno  
                      ch_status) */  
      CUSTNO, INVDATE  
      FROM INVOICE A  
     WHERE CUSTNO = 'DN02'  
           AND STATUS = '90'
```

4 rows
0.74 sec

```
TABLE ACCESS BY ROWID INVOICE  
AND-EQUAL  
INDEX RANGE SCAN CH_STATUS  
INDEX RANGE SCAN CH_CUSTNO
```

```
SELECT CUSTNO, INVDATE  
      FROM INVOICE  
     WHERE CUSTNO = 'DN02'  
           AND INVDATE LIKE '201410%
```

1 row,
0.02 sec

```
TABLE ACCESS BY ROWID INVOICE  
INDEX RANGE SCAN CH_CUSTNO
```

```
SELECT /*+ INDEX(A ch_custno) */  
      CUSTNO, INVDATE  
      FROM INVOICE A  
     WHERE CUSTNO = 'DN02'  
           AND STATUS = '90'
```

4 rows,
0.02 sec

```
TABLE ACCESS BY ROWID INVOICE  
INDEX RANGE SCAN CH_CUSTNO
```

SQL Optimization for DW Index Merge vs. Composite Index Examples

➤ Single Column Index

```
SELECT /*+ AND_EQUAL(A ch_invdte
                     ch_status) */
        CUSTNO, INVDATE
   FROM INVOICE A
  WHERE INVDATE = '941003'
    AND STATUS = '90'
```

13 rows,
0.21 sec

TABLE ACCESS BY ROWID INVOICE
AND-EQUAL
INDEX RANGE SCAN CH_STATUS
INDEX RANGE SCAN CH_INVDTE

```
SELECT CUSTNO, INVDATE
      FROM INVOICE
     WHERE INVDATE LIKE '201501%'
       AND STATUS = '90'
```

TABLE ACCESS BY ROWID INVOICE
INDEX RANGE SCAN CH_STATUS

0 row,
6.05 sec

➤ Composite Index (INVDATE + STATUS)

```
SELECT CUSTNO, INVDATE
      FROM INVOICE A
     WHERE INVDATE = '20150120'
       AND STATUS = '90'
```

13 rows,
sec

```
SELECT CUSTNO, INVDATE
      FROM INVOICE
     WHERE INVDATE LIKE '201501%'
       AND STATUS = '40'
```

rows,
sec

SQL Optimization for DW Composite Index with EQUAL

```
SELECT *  
FROM TAB1  
WHERE COL1 = 'A' AND COL2 = '112'
```

COL1	COL2	ROWID
A	110	10
A	111	11
A	112	5
A	113	18
A	114	54
A	115	23
A	116	29
A	117	25
A	118	26
A	119	30
A	120	19
A	121	32
B	110	41
B	111	45

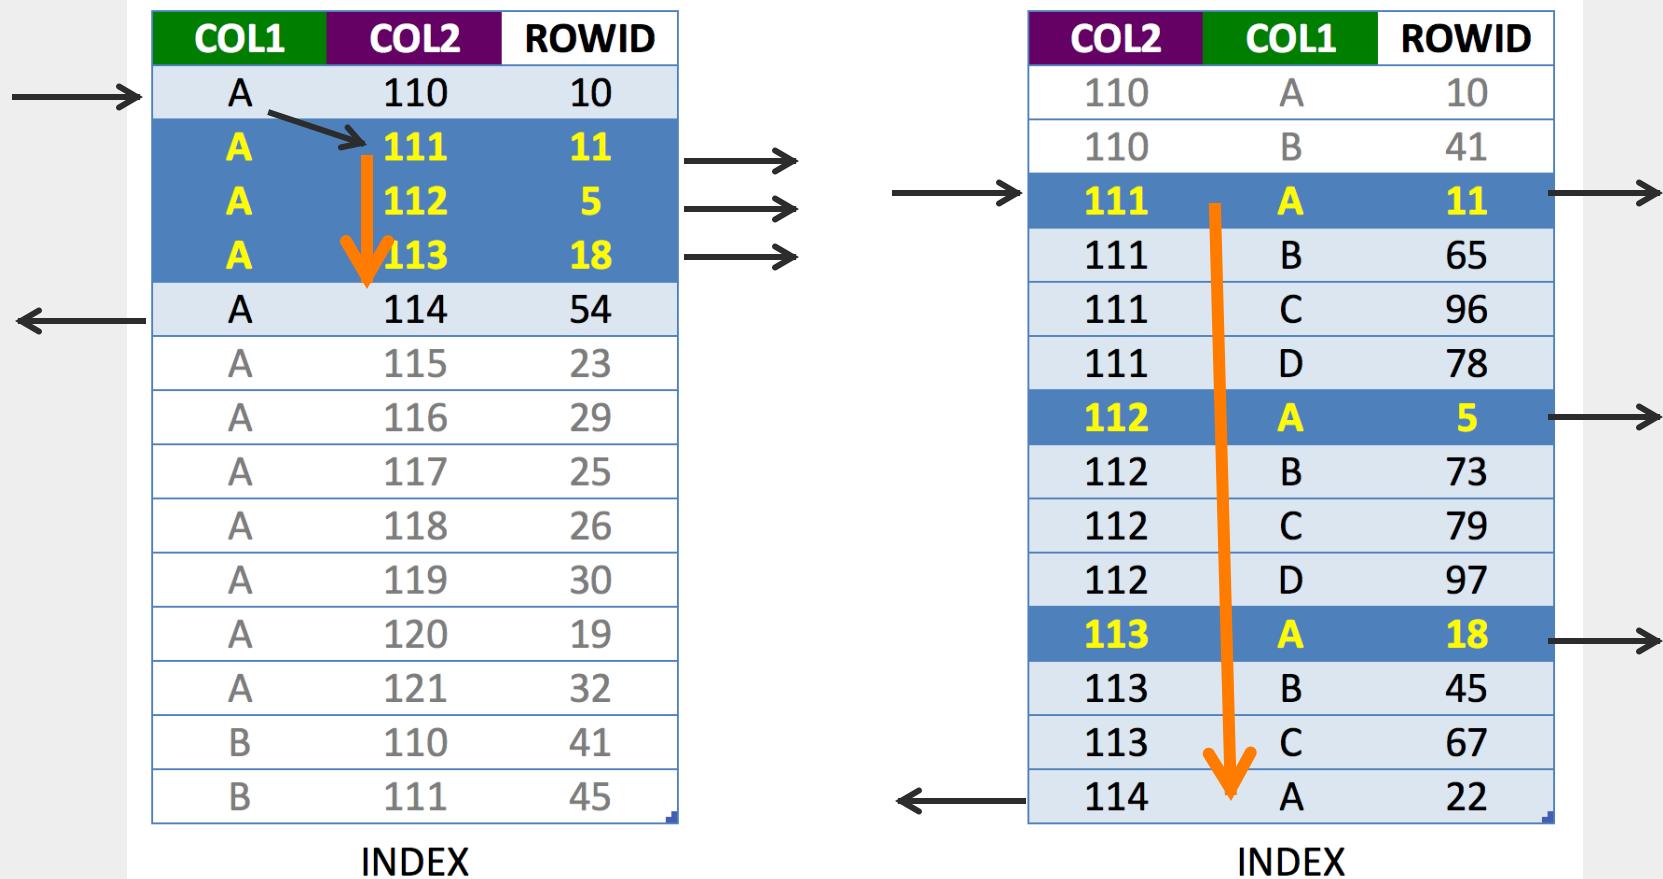
INDEX

COL2	COL1	ROWID
110	A	10
110	B	41
110	C	57
110	D	81
111	A	11
111	B	39
111	C	76
111	D	98
112	A	5
112	B	73
112	C	89
112	D	77
113	A	18
113	B	65

INDEX

SQL Optimization for DW Composite Index with EQUAL (cont'd)

```
SELECT *
  FROM TAB1
 WHERE COL1 = 'A'
   AND COL2 LIKE '111' AND '112'
```



SQL Optimization for DW

Composite Index: BETWEEN vs. IN

```
SELECT * FROM TAB1
WHERE COL1 = 'A'
AND COL2 between '111' AND '112'
```

COL2	COL1	ROWID
110	A	10
110	B	41
111	A	11
111	B	65
111	C	96
111	D	5
112	A	73
112	B	18
112	C	45
114	B	22

INDEX1

TABLE ACCESS BY ROWID TAB1
INDEX RANGE SCAN INDEX1

```
SELECT * FROM TAB1
WHERE COL1 = 'A'
AND COL2 in ('111', '112')
```

COL2	COL1	ROWID
110	A	10
110	B	41
111	A	11
111	B	65
111	C	96
111	D	5
112	A	73
112	B	18
112	C	45
114	B	22

INDEX1

CONCATENATION
TABLE ACCESS BY ROWID TAB1
INDEX RANGE SCAN INDEX1
TABLE ACCESS BY ROWID TAB1
INDEX RANGE SCAN INDEX1

ITMD 526

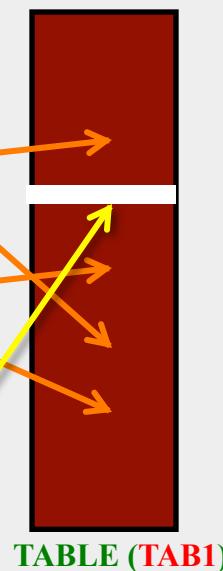
SQL Optimization for DW

Composite Index: Adding a Column

```
SELECT *
  FROM TAB1
 WHERE A = '2' AND C = 51
```

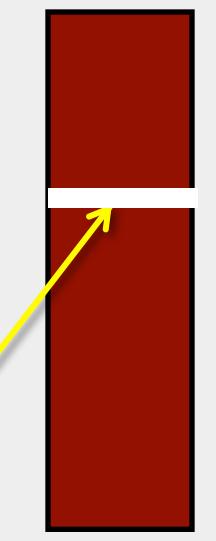
INDEX (A+B)

A	B	ROWID
1	20	10
2	30	11
2	40	47
2	50	54
2	60	58
2	70	61
3	40	21



INDEX (A+B+C)

A	B	C	ROWID
1	20	30	10
2	30	40	11
2	40	45	47
2	50	46	54
2	60	49	58
2	70	51	61
3	40	41	21



SQL Optimization for DW

Composite Index: Adding a Column (cont'd)

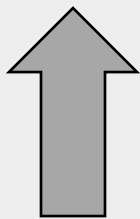
```
SELECT *
  FROM Sales
 WHERE zip = '60201'
   AND dt = '20141005'
   AND st = 'DELIVERED'
```

Index01

zip + dt

Index02

st



Optimizer's Pick

Good

Index01

zip + dt + col

Index02

st



Optimizer's Pick

Bad